

# Bank Marketing Campaign Analysis

## Business Understanding

A Portuguese banking institution recorded data about the potential customers targeted for a phone marketing outreach campaign. This involved one or a series of calls. A more detailed understanding of the results of these campaigns could show which customers are more likely to sign, or show which details about the interactions are most likely to lead to customers signing.

## Data Understanding

The bank provided data on ~45,000 customers, including background information on the client (such as age, job, marital status, etc.) and details about their interactions (such as how frequently and how long ago they were contacted). About 12% of the customers signed. Note that the available data is from Portuguese customers and cultural differences to the US or other countries may mean that the learnings here may not transfer perfectly. Also, a few columns like previous outcome and contact type have a high number of unknown entries, but still contribute significantly to the model. Finally, the model identifies correlations, which do not necessarily imply a causal relationship.

## Definition of columns

- 1 - age (numeric)
- 2 - job : type of job (categorical)
- 3 - marital : marital status (categorical)
- 4 - education (categorical)
- 5 - default: has credit in default? (binary)
- 6 - balance: average yearly balance, in euros (numeric)
- 7 - housing: has housing loan? (binary)
- 8 - loan: has personal loan? (binary)
- 9 - contact: contact communication type (categorical)
- 10 - day: last contact day of the month (numeric)
- 11 - month: last contact month of year (categorical)
- 12 - duration: last contact duration, in seconds (numeric)
- 13 - campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)
- 14 - pdays: number of days since the client was last contacted from a previous campaign (numeric, -1 means client was not previously contacted)
- 15 - previous: number of contacts performed before this campaign and for this client (numeric)
- 16 - poutcome: outcome of the previous marketing campaign (categorical)
- 17 - y - has the client subscribed a term deposit? (binary)

```
In [6]: %pip install pandas numpy scikit-learn matplotlib lightgbm catboost imbalanced-learn xgboo  
  
import pandas as pd  
import numpy as np  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import roc_auc_score  
from sklearn.linear_model import LogisticRegression  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.preprocessing import StandardScaler  
from sklearn.metrics import confusion_matrix  
from imblearn.over_sampling import SMOTE  
import xgboost as xgb  
from xgboost import XGBClassifier  
import lightgbm as lgbm  
from lightgbm import LGBMClassifier  
from itertools import product  
from catboost import CatBoostClassifier, Pool  
import matplotlib.pyplot as plt  
import warnings  
warnings.filterwarnings("ignore", category=DeprecationWarning)
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (2.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (2.0.2)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (1.6.1)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (3.1.0)
Requirement already satisfied: lightgbm in /usr/local/lib/python3.12/dist-packages (4.6.0)
Requirement already satisfied: catboost in /usr/local/lib/python3.12/dist-packages (1.2.8)
Requirement already satisfied: imbalanced-learn in /usr/local/lib/python3.12/dist-packages (0.14.0)
Requirement already satisfied: xgboost in /usr/local/lib/python3.12/dist-packages (3.1.1)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.16.3)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (25.0)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.2.5)
Requirement already satisfied: graphviz in /usr/local/lib/python3.12/dist-packages (from catboost) (0.21)
Requirement already satisfied: plotly in /usr/local/lib/python3.12/dist-packages (from catboost) (5.24.1)
Requirement already satisfied: six in /usr/local/lib/python3.12/dist-packages (from catboost) (1.17.0)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.12/dist-packages (from xgboost) (2.27.3)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.12/dist-packages (from plotly->catboost) (8.5.0)
```

```
In [7]: df=pd.read_csv('https://raw.githubusercontent.com/MattLeRoi/modeling_banking_customers/a5b9df.info()')

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   age         45211 non-null   int64  
 1   job          45211 non-null   object  
 2   marital      45211 non-null   object  
 3   education    45211 non-null   object  
 4   default      45211 non-null   object  
 5   balance      45211 non-null   int64  
 6   housing      45211 non-null   object  
 7   loan          45211 non-null   object  
 8   contact      45211 non-null   object  
 9   day           45211 non-null   int64  
 10  month         45211 non-null   object  
 11  duration     45211 non-null   int64  
 12  campaign     45211 non-null   int64  
 13  pdays        45211 non-null   int64  
 14  previous     45211 non-null   int64  
 15  poutcome     45211 non-null   object  
 16  y             45211 non-null   object  
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```


```
In [8]: df.head()
```

```
Out[8]:   age      job  marital  education  default  balance  housing  loan  contact  day  month  duration  camp
 0   58  management  married  tertiary    no     2143     yes    no  unknown   5  may     261
 1   44  technician  single  secondary   no      29     yes    no  unknown   5  may     151
 2   33 entrepreneur  married  secondary   no       2     yes    yes  unknown   5  may      76
 3   47 blue-collar  married  unknown    no     1506     yes    no  unknown   5  may      92
 4   33      unknown  single  unknown    no       1      no    no  unknown   5  may     198
```

Display histograms and counts of people per category

In [9]:

```
i=1
fig, axes = plt.subplots(4,4, figsize=(18,18))
for col in df.drop('y', axis=1).columns:
    plt.subplot(4,4,i)
    plt.title(col)
    plt.hist(df[col])
    plt.xticks(rotation=45)
    plt.ylabel('Count')
    plt.tight_layout(pad=1.0);
    i+=1
```



Calculate the percentage of people that signed











```
In [10]: y_n_counts = df.y.value_counts()
print("Overall percent signed:",round(y_n_counts[1]/sum(y_n_counts)*100,2),'%')
```

Overall percent signed: 11.7 %

```
/tmp/ipython-input-2652810962.py:2: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
print("Overall percent signed:",round(y_n_counts[1]/sum(y_n_counts)*100,2),'%)
```

Create bar charts with % signed for each column and a reference line for the overall signing rate

```
In [11]: categorical_features = ['day', 'job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',  
temp_y = []  
i=1  
fig, axes = plt.subplots(5,2, figsize=(18,25))  
for cat in categorical_features:  
    temp_y = df.groupby(cat)[‘y’].apply(lambda x: (x == ‘yes’).mean() * 100).sort_values(as  
temp_x = temp_y.index  
plt.subplot(5,2,i)  
plt.bar(temp_x, temp_y)  
plt.title(cat)  
plt.xticks(rotation=45)  
plt.ylabel(‘Percent Signed (%)’)  
plt.axhline(y=11.7, color=‘r’)  
plt.tight_layout(pad=3.0);  
i+=1
```



# Data Preparation

Separate the X and y.

```
In [12]: # The -1 in pdays is a code, not a number, and makes interpretation of the column messy. Since it is a flag for whether the customer was contacted, the pdays column will simply be dropped  
X = df.drop(['y', 'pdays'], axis=1)  
y = [1 if target_y_n == "yes" else 0 for target_y_n in df['y']]
```

Encode the categorical features.

```
In [13]: X_encoded = pd.get_dummies(X, columns=categorical_features)
```

Create the test (15%), validation (10%), and training (75%) sets.

```
In [14]: # Split both encoded and raw (non-encoded) data  
# For models that need encoding: LogisticRegression, RandomForest, LightGBM, and XGBoost  
X_all_training_encoded, X_test_encoded, y_all_training_encoded, y_test_encoded = train_test_split(X_encoded, y, random_state=42, stratify=y)  
X_train_encoded, X_val_encoded, y_train_encoded, y_val_encoded = train_test_split(X_all_training_encoded, y_all_training_encoded, random_state=42, stratify=y)  
  
# For models that can handle categoricals natively: CatBoost  
X_all_training_raw, X_test_raw, y_all_training_raw, y_test_raw = train_test_split(X, y, random_state=42, stratify=y)  
X_train_raw, X_val_raw, y_train_raw, y_val_raw = train_test_split(X_all_training_raw, y_all_training_raw, random_state=42, stratify=y)
```

Scale each column - fit on the training set, then transform val and test sets.

```
In [15]: # Scale only the encoded data (for models that need encoding)  
scaler = StandardScaler()  
X_train_encoded = pd.DataFrame(scaler.fit_transform(X_train_encoded), columns=X_encoded.columns)  
X_val_encoded = pd.DataFrame(scaler.transform(X_val_encoded), columns=X_encoded.columns)  
X_test_encoded = pd.DataFrame(scaler.transform(X_test_encoded), columns=X_encoded.columns)
```


This will loop through a range of ratios to feed into the SMOTE function to analyze the effect of oversampling the churning customers to improve the sensitivity of the model. The results are plotted to visualize the optimum value.

```
In [16]: %time
ratios = np.arange(.2,1,.05)
ROC_AUC_smote = []

for ratio in ratios:
    smote = SMOTE(sampling_strategy=ratio, random_state=42)
    X_train_resampled, y_train_resampled = smote.fit_resample(X_train_encoded, y_train_encoded)

    logreg_resample = LogisticRegression(random_state=42)
    logreg_resample.fit(X_train_resampled, y_train_resampled)
    y_proba = logreg_resample.predict_proba(X_val_encoded)[:, 1]
    score = roc_auc_score(y_val_encoded, y_proba)
    ROC_AUC_smote.append(score)

fig, ax = plt.subplots()
ax.plot(ratios,ROC_AUC_smote,label='ROC_AUC score')
ax.set_ylabel('ROC_AUC score')
ax.set_xlabel('SMOTE ratio')
plt.show();
```



Apply the optimal ratio for SMOTE

```
In [17]: smote = SMOTE(sampling_strategy=0.35, random_state=42)
X_train_encoded, y_train_encoded = smote.fit_resample(X_train_encoded, y_train_encoded)
```

## Modeling

Create a function that takes in a fitted model, calculates the ROC\_AUC score for the validation set, and adds the result to the list of results.

```
In [18]: roc_results = pd.DataFrame(columns=['Model', 'Score'])
imp_factors = pd.DataFrame()
def add_score (roc_results,model_name,model_title, X_val_data, y_val_data):
    y_proba = model_name.predict_proba(X_val_data)[:, 1]
    score = roc_auc_score(y_val_data, y_proba)
    print("ROC-AUC:", score, '\n')
    new_row_data = {'Model':model_title, 'Score':score}
    roc_results.loc[len(roc_results)] = new_row_data
return roc_results
```

## Baseline models

Initialize and fit 5 different models to see which is the most promising for further tuning.

```
In [19]: log_reg = LogisticRegression(random_state=42)
rf=RandomForestClassifier(random_state=42)
xgb_model=XGBClassifier(random_state=42)
lgbm_model=LGBMClassifier(random_state=42,verbose=0)
cat_boost=CatBoostClassifier(random_state=42,verbose=0)

# Train models that need encoding
models_encoded = [('log_reg', log_reg), ('rf', rf), ('xgb_model', xgb_model), ('lgbm_model', lgbm_model),
for title, model in models_encoded:
    print(title)
    model.fit(X_train_encoded, y_train_encoded)
    add_score(roc_results, model, title, X_val_encoded, y_val_encoded)

# For CatBoost, specify categorical features
cat_boost.fit(X_train_raw, y_train_raw, cat_features=categorical_features, verbose=False)
print('cat_boost')
add_score(roc_results, cat_boost, 'cat_boost', X_val_raw, y_val_raw)
```

log\_reg  
ROC-AUC: 0.9136192617178459

rf  
ROC-AUC: 0.9264493563865055

xgb\_model  
ROC-AUC: 0.9327864286022542

lgbm\_model  
ROC-AUC: 0.9330860688378078


cat\_boost  
ROC-AUC: 0.939439324568734

	Model	Score
0	log_reg	0.913619
1	rf	0.926449
2	xgb_model	0.932786
3	lgbm_model	0.933086
4	cat_boost	0.939439

## Identify most important factors from each model and show the top five.


In [20]: # Logistic Regression

```
coefs = pd.Series(log_reg.coef_[0], index=X_train_encoded.columns).sort_values(ascending=False)
coefs.plot(kind='barh');
imp_factors['LogisticRegression']=coefs.sort_values(ascending=False).head(5).index
```



In [21]: # Random Forest

```
importances = pd.Series(rf.feature_importances_, index=X_train_encoded.columns)
importances.sort_values(ascending=False).head(5).sort_values(ascending=True).plot(kind='bar')
imp_factors[ 'RandomForestClassifier']=importances.sort_values(ascending=False).head(5).inde
```




In [22]: # XGBoost

```
importance = xgb_model.get_booster().get_score(importance_type='gain')
importance_df = pd.DataFrame({'importance':list(importance.values())}, index=importance.keys)

importances=importance_df.sort_values(by='importance', ascending=True).tail(5)
importances.plot(kind='barh');

imp_factors['XGBClassifier']=importances.sort_values(by='importance', ascending=False).head(5)
```




In [23]: # Light GBM

```
importance = lgbm_model.booster_.feature_importance(importance_type='gain')
feature_names = lgbm_model.feature_name_

importance_df = pd.DataFrame({'importance':importance}, index=feature_names)

importances=importance_df.sort_values(by='importance', ascending=True).tail(5)
importances.plot(kind='barh');
imp_factors['LGBMClassifier']=importances.sort_values(by='importance', ascending=False).head(5)
```




In [24]: # Cat Boost

```
importance = cat_boost.get_feature_importance()
feature_names = X_train_raw.columns

importance_df = pd.DataFrame({'importance':importance}, index=feature_names)

importances=importance_df.sort_values(by='importance', ascending=True).tail(5)
importances.plot(kind='barh');
imp_factors['CatBoostClassifier']=importances.sort_values(by='importance', ascending=False)
```



These are the top five factors for each model. I created a composite score to find out, across all of the models, which factors showed up as consistently important (for each model, 1st place is 5 points, 2nd place is 4 points, etc.).

```
In [25]: factor_importance_results = pd.DataFrame(columns=['score'], index=X_train_encoded.columns)

for cat in X_train_encoded.columns:
    tot_factor_weight=0
    for col in imp_factors:
        ind_factor_weight = 5*imp_factors[imp_factors[col].str.contains(cat, case=False, na=False)]
        if ind_factor_weight > 0:
            tot_factor_weight += ind_factor_weight
    if tot_factor_weight > 0:
        factor_importance_results.at[cat, 'score'] = tot_factor_weight[0]
    else: factor_importance_results.at[cat, 'score'] = 0
display(imp_factors)
factor_importance_results.sort_values(by='score', ascending=False).head(5)
```

	LogisticRegression	RandomForestClassifier	XGBClassifier	LGBMClassifier	CatBoostClassifier
0	duration		duration	poutcome_success	duration
1	poutcome_success		balance	loan_yes	housing_no
2	contact_cellular		age	contact_unknown	poutcome_success
3	month_jun		campaign	housing_yes	campaign
4	month_mar		housing_no	housing_no	contact_unknown
					poutcome

```
Out[25]:
```

	score
duration	20
poutcome_success	12
housing_no	6
balance	4
contact_unknown	4

## Tuning

For the CatBoostClassifier, several parameters (learning\_rates, depths, l2\_leaf\_regs, and class\_weights) are varied to see which yields the highest ROC\_AUC score on the validation set.

```
In [26]: # This is the best result from previous runs. To save run time it is manually entered here.
# takes ~20 minutes, uncomment the next cell and it will overwrite this result.
results_df = pd.DataFrame([(8, 0.07, 1, 3, 0.940960)], columns=['depth', 'learning_rate',
```

```
In [27]: train_pool = Pool(X_train_raw, y_train_raw, cat_features=categorical_features)
val_pool = Pool(X_val_raw, y_val_raw, cat_features=categorical_features)
```

In [28]: %time

```
# # Uncomment this section to run the full parameter set (takes ~20 minutes)
# depths = [6,8,10]
# Learning_rates = [0.03, .05, .07]
# L2_Leaf_regs = [1,3]
# class_weights = [1,3]
# total_runs = product(Learning_rates, depths, L2_Leaf_regs, class_weights)
# num_runs = len(depths)*len(Learning_rates)*len(L2_Leaf_regs)*len(class_weights)

# results = []

# i=1
# for lr, d, l2, cw in total_runs:
#     model = CatBoostClassifier(
#         learning_rate=lr,
#         depth=d,
#         l2_leaf_reg=l2,
#         iterations=1000,
#         early_stopping_rounds=50,
#         use_best_model=True,
#         verbose=False,
#         random_seed=42,
#         class_weights=[1,cw]
#     )
#     model.fit(train_pool, eval_set=val_pool, early_stopping_rounds=50)
#     y_proba = model.predict_proba(X_val_raw)[:, 1]
#     score = roc_auc_score(y_val_raw, y_proba)
#     results.append((d, lr, l2, cw, score))
#     print(i, ' of ', num_runs)
#     i+=1

# results_df = pd.DataFrame(results, columns=['depth', 'Learning_rate', 'L2_Leaf_reg', 'class_weights', 'score'])
# results_df=results_df.sort_values('AUC', ascending=False)
# display(results_df.head())
```

CPU times: user 3 µs, sys: 0 ns, total: 3 µs  
Wall time: 5.48 µs

## Evaluation

Run the final tuned model with the validation data for comparison to the baseline models.

In [29]:

```
%time
best_params = {
    'depth': results_df['depth'].iloc[0],
    'learning_rate': results_df['learning_rate'].iloc[0],
    'iterations': 1000,
    'l2_leaf_reg': results_df['l2_leaf_reg'].iloc[0],
    'loss_function': 'Logloss',
    'eval_metric': 'AUC',
    'random_seed': 42,
    'early_stopping_rounds': 50,
    'use_best_model': True,
    'verbose': False,
    'class_weights': [1, results_df['class_weights'].iloc[0]]
}

final_model = CatBoostClassifier(**best_params);
final_model.fit(train_pool, eval_set=val_pool, early_stopping_rounds=50);
add_score(roc_results, final_model, 'final_model', X_val_raw, y_val_raw)
```

ROC-AUC: 0.9403338315894371

CPU times: user 1min, sys: 3.38 s, total: 1min 4s  
Wall time: 37.5 s

Out[29]:

	Model	Score
0	log_reg	0.913619
1	rf	0.926449
2	xgb_model	0.932786
3	lgbm_model	0.933086
4	cat_boost	0.939439
5	final_model	0.940334

Run the final model on the test data.

In [30]:

```
y_pred = final_model.predict(X_test_raw)
y_proba = final_model.predict_proba(X_test_raw)[:, 1]

print(confusion_matrix(y_test_raw,y_pred))
print("Final ROC-AUC score:", round(roc_auc_score(y_test_raw,y_proba),4))
```

[[5432 538]  
 [ 180 632]]

Final ROC-AUC score: 0.9357

## Next steps:

1. Collect more detailed data - filling in all of the missing data may lead to further insights.
2. Run a controlled experiment - design an experiment to determine whether call duration and number of contacts are causal or correlative relationships.
3. The data varied by month, both in volume and in the success rate. This may indicate seasonality or it may have been due to the timing of the marketing campaign, which I don not have access to. It is worth further

investigation